

# **SANDIA REPORT**

SAND2014-18874

Unlimited Release

Printed October 2014

## **MueLu User's Guide 1.0**

Andrey Prokopenko, Jonathan J. Hu, Tobias A. Wiesner, Christopher M. Siefert,  
Raymond S. Tuminaro

Prepared by  
Sandia National Laboratories  
Albuquerque, New Mexico 87185 and Livermore, California 94550

Sandia National Laboratories is a multi-program laboratory managed and operated by Sandia Corporation, a wholly owned subsidiary of Lockheed Martin Corporation, for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-AC04-94AL85000.

Approved for public release; further dissemination unlimited.



**Sandia National Laboratories**

Issued by Sandia National Laboratories, operated for the United States Department of Energy by Sandia Corporation.

**NOTICE:** This report was prepared as an account of work sponsored by an agency of the United States Government. Neither the United States Government, nor any agency thereof, nor any of their employees, nor any of their contractors, subcontractors, or their employees, make any warranty, express or implied, or assume any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represent that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise, does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States Government, any agency thereof, or any of their contractors or subcontractors. The views and opinions expressed herein do not necessarily state or reflect those of the United States Government, any agency thereof, or any of their contractors.

Printed in the United States of America. This report has been reproduced directly from the best available copy.

Available to DOE and DOE contractors from  
U.S. Department of Energy  
Office of Scientific and Technical Information  
P.O. Box 62  
Oak Ridge, TN 37831

Telephone: (865) 576-8401  
Facsimile: (865) 576-5728  
E-Mail: [reports@adonis.osti.gov](mailto:reports@adonis.osti.gov)  
Online ordering: <http://www.osti.gov/bridge>

Available to the public from  
U.S. Department of Commerce  
National Technical Information Service  
5285 Port Royal Rd  
Springfield, VA 22161

Telephone: (800) 553-6847  
Facsimile: (703) 605-6900  
E-Mail: [orders@ntis.fedworld.gov](mailto:orders@ntis.fedworld.gov)  
Online ordering: <http://www.ntis.gov/help/ordermethods.asp?loc=7-4-0#online>



## MueLu User's Guide 1.0

Andrey Prokopenko  
Scalable Algorithms  
Sandia National Laboratories  
Mailstop 1318  
P.O. Box 5800  
Albuquerque, NM 87185-1318  
aprokop@sandia.gov

Tobias Wiesner  
Institute for Computational Mechanics  
Technische Universität München  
Boltzmanstraße 15  
85747 Garching, Germany  
wiesner@lnm.mw.tum.de

Jonathan J. Hu  
Scalable Algorithms  
Sandia National Laboratories  
Mailstop 9159  
P.O. Box 0969  
Livermore, CA 94551-0969  
jhu@sandia.gov

Christopher M. Siefert  
Computational Multiphysics  
Sandia National Laboratories  
Mailstop 1322  
P.O. Box 5800  
Albuquerque, NM 87185-1322  
csiefer@sandia.gov

Raymond S. Tuminaro  
Computational Mathematics  
Sandia National Laboratories  
Mailstop 9159  
P.O. Box 0969  
Livermore, CA 94551-0969  
rstumin@sandia.gov

## Abstract

This is the official user guide for MUELU multigrid library in Trilinos version 12.6. This guide provides an overview of MUELU, its capabilities, and instructions for new users who want to start using MUELU with a minimum of effort. Detailed information is given on how to drive MUELU through its XML interface. Links to more advanced use cases are given. This guide gives information on how to achieve good parallel performance, as well as how to introduce new algorithms. Finally, readers will find a comprehensive listing of available MUELU options. *Any options not documented in this manual should be considered strictly experimental.*

# Acknowledgment

Many people have helped develop MUELU and/or provided valuable feedback, and we would like to acknowledge their contributions here: Tom Benson, Julian Cortial, Eric Cyr, Stefan Domino, Travis Fisher, Jeremie Gaidamour, Axel Gerstenberger, Chetan Jhurani, Mark Hoemmen, Paul Lin, Eric Phipps, Siva Rajamanickam, Nico Schlömer, and Paul Tsuji.



# Contents

<b>1</b>	<b>Introduction</b>	<b>13</b>
<b>2</b>	<b>Multigrid background</b>	<b>15</b>
<b>3</b>	<b>Getting Started</b>	<b>17</b>
3.1	Overview of MUELU .....	17
3.2	Configuration and Build .....	18
3.2.1	Dependencies .....	18
3.2.2	Configuration .....	20
3.3	Examples in code .....	20
3.3.1	MUELU as a preconditioner within BELOS .....	21
3.3.2	MUELU as a preconditioner for AZTECOO .....	23
3.3.3	Further remarks .....	24
<b>4</b>	<b>Performance tips</b>	<b>25</b>
<b>5</b>	<b>MUELU options</b>	<b>27</b>
5.1	Using parameters on individual levels .....	27
5.2	Parameter validation .....	27
5.3	General options .....	28
5.4	Smoothing and coarse solver options .....	29
5.5	Aggregation options .....	32
5.6	Rebalancing options .....	33
5.7	Multigrid algorithm options .....	34

5.8	Reuse options .....	36
5.9	Miscellaneous options .....	37
<b>6</b>	<b>MUEMEX: The MATLAB Interface for MUELU</b>	<b>39</b>
6.1	Cmake Configure and Make .....	39
6.1.1	BLAS & LAPACK Option #1: Static Builds .....	40
6.1.2	BLAS & LAPACK Option #2: LD_PRELOAD .....	41
6.1.3	Running MATLAB .....	41
6.2	Using MUEMEX .....	42
6.2.1	Setup Mode .....	42
6.2.2	Solve Mode .....	43
6.2.3	Cleanup Mode .....	43
6.2.4	Status Mode .....	43
6.2.5	Get Mode .....	44
6.2.6	Tips and Tricks .....	44
	<b>References</b>	<b>45</b>
 <b>Appendix</b>		
<b>A</b>	<b>Copyright and License</b>	<b>47</b>
<b>B</b>	<b>ML compatibility</b>	<b>49</b>
B.1	Usage of ML parameter lists with MUELU .....	49
B.1.1	MLParameterListInterpreter .....	49
B.1.2	ML2MueLuParameterTranslator .....	50
B.2	Compatible ML parameters .....	51
B.2.1	General ML options .....	51



B.2.2	Smoothing and coarse solver options .....	52
B.2.3	Transfer operator options .....	53
B.2.4	Rebalancing options .....	53

# List of Figures

# List of Tables

3.1	MUELU's required and optional dependencies, subdivided by whether a dependency is that of the MUELU library itself ( <i>Library</i> ), or of some MUELU test ( <i>Testing</i> ). . . . .	19
5.1	Verbosity levels. . . . .	28
5.2	Supported problem types ("—" means not set). . . . .	28
5.3	Commonly used smoothers provided by IFPACK/IFPACK2. Note that these smoothers can also be used as coarse grid solvers. . . . .	30
5.4	Commonly used direct solvers provided by AMESOS/AMESOS2 . . . . .	30
5.5	Available coarsening schemes. . . . .	32
5.6	Available multigrid algorithms for generating grid transfer matrices. . . . .	35
5.7	Available coarsening schemes. . . . .	36



# Chapter 1

## Introduction

This guide gives an overview of MUELU’s capabilities. If you are looking for a tutorial, please refer to the MUELU tutorial in `muelu/doc/Tutorial` (see also [17]). New users should start with §3. It strives to give the new user all the information he/she might need to begin using MUELU quickly. Users interested in performance, especially in parallel context, should refer to §4. Users looking for a particular option should consult §5, containing a complete set of supported options in MUELU.

If you find any errors or omissions in this guide, have comments or suggestions, or would like to contribute to MUELU development, please contact the MUELU [users list](#), or [developers list](#).



# Chapter 2

## Multigrid background

Here we provide a brief multigrid introduction (see [5] or [14] for more information). A multigrid solver tries to approximate the original problem of interest with a sequence of smaller (*coarser*) problems. The solutions from the coarser problems are combined in order to accelerate convergence of the original (*fine*) problem on the finest grid. A simple multilevel iteration is illustrated in Algorithm 1.

---

**Algorithm 1** V-cycle multigrid with  $N$  levels to solve  $Ax = b$ .

---

```
 $A_0 = A$ 
function MULTILEVEL( $A_k, b, u, k$ )
  // Solve  $A_k u = b$  ( $k$  is current grid level)
   $u = S_m^1(A_k, b, u)$ 
  if ( $k \neq N - 1$ ) then
     $P_k = \text{determine\_interpolant}(A_k)$ 
     $R_k = \text{determine\_restrictor}(A_k)$ 
     $\hat{r}_{k+1} = R_k(b - A_k u)$ 
     $A_{k+1} = R_k A_k P_k$ 
     $v = 0$ 
    MULTILEVEL( $\hat{A}_{k+1}, \hat{r}_{k+1}, v, k + 1$ )
     $u = u + P_k v$ 
     $u = S_m^2(A_k, b, u)$ 
  end if
end function
```

---

In the multigrid iteration in Algorithm 1, the  $S_m^1()$ 's and  $S_m^2()$ 's are called *pre-smoothers* and *post-smoothers*. They are approximate solvers (e.g., symmetric Gauss-Seidel), with the subscript  $m$  denoting the number of applications of the approximate solution method. The purpose of a smoother is to quickly reduce certain error modes in the approximate solution on a level  $k$ . For symmetric problems, the pre- and post-smoothers should be chosen to maintain symmetry (e.g., forward Gauss-Seidel for the pre-smoother and backward Gauss-Seidel for the post-smoother). The  $P_k$ 's are *interpolation* matrices that transfer solutions from coarse levels to finer levels. The  $R_k$ 's are *restriction* matrices that restrict a fine level solution to a coarser level. In a geometric multigrid,  $P_k$ 's and  $R_k$ 's are determined by the application, whereas in an algebraic multigrid they are automatically generated. For symmetric problems, typically  $R_k = P_k^T$ . For nonsymmetric problems, this is not necessarily true. The  $A_k$ 's are the coarse level problems, and are generated through

a Galerkin (triple matrix) product.

Please note that the algebraic multigrid algorithms implemented in MUELU generate the grid transfers  $P_k$  automatically and the coarse problems  $A_k$  via a sparse triple matrix product. TRILINOS provides a wide selection of smoothers and direct solvers for use in MUELU through the IFPACK, IFPACK2, AMESOS, and AMESOS2 packages (see §5).



# Chapter 3

## Getting Started

This section is meant to get you using MUELU as quickly as possible. §3 gives a summary of MUELU’s design. §3.2 lists MUELU’s dependencies on other TRILINOS libraries and provides a sample cmake configuration line. Finally, code examples using the XML interface are given in §3.3.

### 3.1 Overview of MUELU

MUELU is an extensible algebraic multigrid (AMG) library that is part of the TRILINOS project. MUELU works with EPETRA (32-bit version <sup>1</sup>) and TPETRA matrix types. The library is written in C++ and allows for different ordinal (index) and scalar types. MUELU is designed to be efficient on many different computer architectures, from workstations to supercomputers, relying on “MPI+X” principle, where “X” can be threading or CUDA.

MUELU provides a number of different multigrid algorithms:

1. smoothed aggregation AMG (for Poisson-like and elasticity problems);
2. Petrov-Galerkin aggregation AMG (for convection-diffusion problems);
3. energy-minimizing AMG;
4. aggregation-based AMG for problems arising from the eddy current formulation of Maxwell’s equations.

MUELU’s software design allows for the rapid introduction of new multigrid algorithms. The most important features of MUELU can be summarized as:

#### Easy-to-use interface

MUELU has a user-friendly parameter input deck which covers most important use cases. Reasonable defaults are provided for common problem types (see Table 5.2).

---

<sup>1</sup>Support for the Epetra 64-bit version is planned.

## **Modern object-oriented software architecture**

MUELU is written completely in C++ as a modular object-oriented multigrid framework, which provides flexibility to combine and reuse existing components to develop novel multigrid methods.

## **Extensibility**

Due to its flexible design, MUELU is an excellent toolkit for research on novel multigrid concepts. Experienced multigrid users have full access to the underlying framework through an advanced XML based interface. Expert users may use and extend the C++ API directly.

## **Integration with TRILINOS library**

As a package of TRILINOS, MUELU is well integrated into the TRILINOS environment. MUELU can be used with either the TPETRA or EPETRA (32-bit) linear algebra stack. It is templated on the local index, global index, scalar, and compute node types. This makes MUELU ready for future developments.

## **Broad range of supported platforms**

MUELU runs on wide variety of architectures, from desktop workstations to parallel Linux clusters and supercomputers ( [9]).

## **Open source**

MUELU is freely available through a simplified BSD license (see Appendix [A](#)).

# **3.2 Configuration and Build**

MUELU has been compiled successfully under Linux with the following C++ compilers: GNU versions 4.1 and later, Intel versions 12.1/13.1, and clang versions 3.2 and later. In the future, we recommend using compilers supporting C++11 standard.

## **3.2.1 Dependencies**

### **Required Dependencies**

MUELU requires that TEUCHOS and either EPETRA/IFPACK or TPETRA/IFPACK2 are enabled.

### **Recommended Dependencies**

We strongly recommend that you enable the following TRILINOS libraries along with MUELU:

- EPETRA stack: AZTECOO, EPETRA, AMESOS, IFPACK, ISORROPIA, GALERI, ZOLTAN;

- TPETRA stack: AMESOS2, BELOS, GALERI, IFPACK2, TPETRA, ZOLTAN2.

## Tutorial Dependencies

In order to run the MUELU Tutorial [17] located in `muelu/doc/Tutorial`, MUELU must be configured with the following dependencies enabled:

AZTECOO, AMESOS, AMESOS2, BELOS, EPETRA, IFPACK, IFPACK2, ISORROPIA, GALERI, TPETRA, ZOLTAN, ZOLTAN2.

☛ Note that the MUELU tutorial [17] comes with a VirtualBox image with a pre-installed Linux and TRILINOS. In this way, a user can immediately begin experimenting with MUELU without having to install the TRILINOS libraries. Therefore, it is an ideal starting point to get in touch with MUELU.

## Complete List of Direct Dependencies

Dependency	Required		Optional	
	Library	Testing	Library	Testing
AMESOS			×	×
AMESOS2			×	×
AZTECOO				×
BELOS				×
EPETRA			×	×
IFPACK			×	×
IFPACK2			×	×
ISORROPIA			×	×
GALERI				×
KOKKOSCLASSIC			×	
TEUCHOS	×	×		
TPETRA			×	×
XPETRA	×	×		
ZOLTAN			×	×
ZOLTAN2			×	×
Boost			×	
BLAS	×	×		
LAPACK	×	×		
MPI			×	×

**Table 3.1.** MUELU’s required and optional dependencies, subdivided by whether a dependency is that of the MUELU library itself (*Library*), or of some MUELU test (*Testing*).

Table 3.1 lists the dependencies of MUELU, both required and optional. If an optional dependency is not present, the tests requiring that dependency are not built.

☛ AMESOS/AMESOS2 are necessary if one wants to use a sparse direct solve on the coarsest level. ZOLTAN/ZOLTAN2 are necessary if one wants to use matrix rebalancing in parallel runs (see §4). AZTECOO/BELOS are necessary if one wants to test MUELU as a preconditioner instead of a solver.

☛ MUELU has also been successfully tested with SuperLU 4.1 and SuperLU 4.2.

☛ Some packages that MUELU depends on may come with additional requirements for third party libraries, which are not listed here as explicit dependencies of MUELU. It is highly recommended to install ParMetis 3.1.1 or newer for ZOLTAN, ParMetis 4.0.x for ZOLTAN2, and SuperLU 4.1 or newer for AMESOS/AMESOS2.

### 3.2.2 Configuration

The preferred way to configure and build MUELU is to do that outside of the source directory. Here we provide a sample configure script that will enable MUELU and all of its optional dependencies:

```
export TRILINOS_HOME=/path/to/your/Trilinos/source/directory
cmake \
  -D BUILD_SHARED_LIBS:BOOL=ON \
  -D CMAKE_BUILD_TYPE:STRING="RELEASE" \
  -D CMAKE_CXX_FLAGS:STRING="-g" \
  -D Trilinos_ENABLE_EXPLICIT_INSTANTIATION:BOOL=ON \
  -D Trilinos_ENABLE_TESTS:BOOL=OFF \
  -D Trilinos_ENABLE_EXAMPLES:BOOL=OFF \
  -D Trilinos_ENABLE_MueLu:BOOL=ON \
  -D MueLu_ENABLE_TESTS:STRING=ON \
  -D MueLu_ENABLE_EXAMPLES:STRING=ON \
  -D TPL_ENABLE_BLAS:BOOL=ON \
  -D TPL_ENABLE_MPI:BOOL=ON \
  ${TRILINOS_HOME}
```

More configure examples can be found in Trilinos/sampleScripts. For more information on configuring, see the TRILINOS Cmake Quickstart guide [1].

## 3.3 Examples in code

The most commonly used scenario involving MUELU is using a multigrid preconditioner inside an iterative linear solver. In TRILINOS, a user has a choice between EPETRA and TPETRA

for the underlying linear algebra library. Important Krylov subspace methods (such as preconditioned CG and GMRES) are provided in TRILINOS packages AZTECOO (EPETRA) and BELOS (EPETRA/TPETRA).

At this point, we assume that the reader is comfortable with TEUCHOS referenced-counted pointers (RCPs) for memory management (an introduction to RCPs can be found in [3]) and the `Teuchos::ParameterList` class [13].

### 3.3.1 MUELU as a preconditioner within BELOS

The following code shows the basic steps of how to use a MUELU multigrid preconditioner with TPETRA linear algebra library and with a linear solver from BELOS. To keep the example short and clear, we skip the template parameters and focus on the algorithmic outline of setting up a linear solver. For further details, a user may refer to the examples and test directories.

First, we create the MUELU multigrid preconditioner. It can be done in a few ways. For instance, multigrid parameters can be read from an XML file (e.g., *mueluOptions.xml* in the example below).

```
Teuchos::RCP<Tpetra::CrsMatrix<> > A;
// create A here ...
std::string optionsFile = "mueluOptions.xml";
Teuchos::RCP<MueLu::TpetraOperator> mueluPreconditioner =
    MueLu::CreateTpetraPreconditioner(A, optionsFile);
```

The XML file contains multigrid options. A typical file with MUELU parameters looks like the following.

```
<ParameterList name="MueLu">

  <Parameter name="verbosity" type="string" value="low"/>

  <Parameter name="max levels" type="int" value="3"/>
  <Parameter name="coarse: max size" type="int" value="10"/>

  <Parameter name="multigrid algorithm" type="string" value="sa"/>

  <!-- Damped Jacobi smoothing -->
  <Parameter name="smoother: type" type="string" value="RELAXATION"/>
  <ParameterList name="smoother: params">
    <Parameter name="relaxation: type" type="string" value="Jacobi"/>
    <Parameter name="relaxation: sweeps" type="int" value="1"/>
    <Parameter name="relaxation: damping factor" type="double" value="0.9"/>
  </ParameterList>

  <!-- Aggregation -->
```

```

<Parameter name="aggregation: type" type="string" value="uncoupled"/>
<Parameter name="aggregation: min agg size" type="int" value="3"/>
<Parameter name="aggregation: max agg size" type="int" value="9"/>

</ParameterList>

```

It defines a three level smoothed aggregation multigrid algorithm. The aggregation size is between three and nine(2D)/27(3D) nodes. One sweep with a damped Jacobi method is used as a level smoother. By default, a direct solver is applied on the coarsest level. A complete list of available parameters and valid parameter choices is given in §5 of this User's Guide.

Users can also construct a multigrid preconditioner using a provided ParameterList without accessing any files in the following manner.

```

Teuchos::RCP<Tpetra::CrsMatrix<> > A;
// create A here ...
Teuchos::ParameterList paramList;
paramList.set("verbosity", "low");
paramList.set("max levels", 3);
paramList.set("coarse: max size", 10);
paramList.set("multigrid algorithm", "sa");
// ...
Teuchos::RCP<MueLu::TpetraOperator> mueLuPreconditioner =
    MueLu::CreateTpetraPreconditioner(A, paramList);

```

Besides the linear operator  $A$ , we also need an initial guess vector for the solution  $X$  and a right hand side vector  $B$  for solving a linear system.

```

Teuchos::RCP<const Tpetra::Map<> > map = A->getDomainMap();

// Create initial vectors
Teuchos::RCP<Tpetra::MultiVector<> > B, X;
X = Teuchos::rcp( new Tpetra::MultiVector<>(map,numrhs) );
Belos::MultiVecTraits<>::MvRandom( *X );
B = Teuchos::rcp( new Tpetra::MultiVector<>(map,numrhs) );
Belos::OperatorTraits<>::Apply( *A, *X, *B );
Belos::MultiVecTraits<>::MvInit( *X, 0.0 );

```

To generate a dummy example, the above code first declares two vectors. Then, a right hand side vector is calculated as the matrix-vector product of a random vector with the operator  $A$ . Finally, an initial guess is initialized with zeros.

Then, one can define a `Belos::LinearProblem` object where the `mueLuPreconditioner` is used for left preconditioning

```

Belos::LinearProblem<> problem( A, X, B );
problem->setLeftPrec(mueLuPreconditioner);

```

```
bool set = problem.setProblem();
```

Next, we set up a BELOS solver using some basic parameters

```
Teuchos::ParameterList belosList;  
belosList.set( "Block Size", 1 );  
belosList.set( "Maximum Iterations", 100 );  
belosList.set( "Convergence Tolerance", 1e-10 );  
belosList.set( "Output Frequency", 1 );  
belosList.set( "Verbosity", Belos::TimingDetails + Belos::FinalSummary );  
  
Belos::BlockCGSolMgr<> solver( rcp(&problem,false), rcp(&belosList,false) );
```

Finally, we solve the system.

```
Belos::ReturnType ret = solver.solve();
```

### 3.3.2 MUELU as a preconditioner for AZTECOO

For EPETRA, users have two library options: BELOS (recommended) and AZTECOO. AZTECOO and BELOS both provide fast and mature implementations of common iterative Krylov linear solvers. BELOS has additional capabilities, such as Krylov subspace recycling and “tall skinny QR”.

Constructing a MUELU preconditioner for Epetra operators is done in a similar manner to Tpetra.

```
Teuchos::RCP<Epetra_CrsMatrix> A;  
// create A here ...  
Teuchos::RCP<MueLu::EpetraOperator> mueLuPreconditioner;  
std::string optionsFile = "mueluOptions.xml";  
mueLuPreconditioner = MueLu::CreateEpetraPreconditioner(A, optionsFile);
```

MUELU parameters are generally Epetra/Tpetra agnostic, thus the XML parameter file could be the same as §3.3.1.

Furthermore, we assume that a right hand side vector and a solution vector with the initial guess are defined.

```
Teuchos::RCP<const Epetra_Map> map = A->DomainMap();  
Teuchos::RCP<Epetra_Vector> B = Teuchos::rcp(new Epetra_Vector(map));  
Teuchos::RCP<Epetra_Vector> X = Teuchos::rcp(new Epetra_Vector(map));  
X->PutScalar(0.0);
```

Then, an Epetra.LinearProblem can be defined.

```
Epetra_LinearProblem epetraProblem(A.get(), X.get(), B.get());
```

The following code constructs an AZTECOO CG solver.

```
Aztec00 aztecSolver(epetraProblem);  
aztecSolver.SetAztecOption(AZ_solver, AZ_cg);  
aztecSolver.SetPrecOperator(mueLuPreconditioner.get());
```

Finally, the linear system is solved.

```
int maxIts = 100;  
double tol = 1e-10;  
aztecSolver.Iterate(maxIts, tol);
```

### 3.3.3 Further remarks

This section is only meant to give a brief introduction on how to use MUELU as a preconditioner within the TRILINOS packages for iterative solvers. There are other, more complicated, ways to use MUELU as a preconditioner for BELOS and AZTECOO through the XPETRA interface. Of course, MUELU can also work as standalone multigrid solver. For more information on these topics, the reader may refer to the examples and tests in the MUELU source directory (Trilinos/packages/muelu), as well as to the MUELU tutorial [17]. For in-depth details of MUELU applied to multiphysics problems, please see [16].



# Chapter 4

## Performance tips

In practice, it can be very challenging to find an appropriate set of multigrid parameters for a specific problem, especially if few details are known about the underlying linear system. In this Chapter, we provide some advice for improving multigrid performance.

☛ For optimizing multigrid parameters, it is highly recommended to set the verbosity to `high` or `extreme` for MUELU to output more information (e.g., for the effect of the chosen parameters to the aggregation and coarsening process).

Some general advice:

- Choose appropriate iterative linear solver (e.g., GMRES for non-symmetric problems).
- Start with the recommended settings for particular problem types. See Table 5.2.
- Choose reasonable basic multigrid parameters (see §5.3), including maximum number of multigrid levels (`max_levels`) and maximum allowed coarse size of the problem (`coarse_max_size`). Take fine level problem size and sparsity pattern into account for a reasonable choice of these parameters.
- Select an appropriate transfer operator strategy (see §5.7). For symmetric problems, you should start with smoothed aggregation multigrid. For non-symmetric problems, a Petrov-Galerkin smoothed aggregation method is a good starting point, though smoothed aggregation may also perform well.
- Enable implicit restrictor construction (`transpose: use_implicit`) for symmetric problems.
- Find good level smoothers (see §5.4). If a problem is symmetric positive definite, choose a smoother with a matrix-vector computational kernel, such as the Chebyshev polynomial smoother. If you are using relaxation smoothers, we recommend starting with optimizing the damping parameter. Once you have found a good damping parameter for your problem, you can increase the number of smoothing iterations.
- Adjust aggregation parameters if you experience bad coarsening ratios (see §5.5). Particularly, try adjusting the minimum (`aggregation: min_agg_size`) and maximum (`aggregation: max_agg_size`) aggregation parameters. For a 2D (3D) isotropic problem on a regular mesh, the aggregate size should be about 9 (27) nodes per aggregate.

- Replace a direct solver with an iterative method (`coarse: type`) if your coarse level solution becomes too expensive (see §5.4).

Some advice for parallel runs include:

1. Enable matrix rebalancing when running in parallel (`repartition: enable`).
2. Use smoothers invariant to the number of processors, such as polynomial smoothing, if possible.
3. Use uncoupled aggregation instead of coupled, as later requires significantly more communication.
4. Adjust rebalancing parameters (see §5.6). Try choosing rebalancing parameters so that you end up with one processor on the coarsest level for the direct solver (this avoids additional communication).
5. Enable implicit rebalancing of prolongators and restrictors (`repartition: rebalance P and R`).

# Chapter 5

## MUELU options

In this section, we report the complete list of MUELU input parameters. It is important to notice, however, that MUELU relies on other TRILINOS packages to provide support for some of its algorithms. For instance, IFPACK/IFPACK2 provide standard smoothers like Jacobi, Gauss-Seidel or Chebyshev, while AMESOS/AMESOS2 provide access to direct solvers. The parameters affecting the behavior of the algorithms in those packages are simply passed by MUELU to a routine from the corresponding library. Please consult corresponding packages' documentation for a full list of supported algorithms and corresponding parameters.

### 5.1 Using parameters on individual levels

Some of the parameters that affect the preconditioner can in principle be different from level to level. By default, parameters affect all levels in a multigrid hierarchy.

The settings on a particular levels can be changed by using level sublists. A level sublist is a `ParameterList` sublist with the name “level XX”, where XX is the level number. The parameter names in the sublist do not require any modifications. For example, the following fragment of code

```
<ParameterList name="level 2">
  <Parameter name="smoother: type" type="string" value="CHEBYSHEV"/>
</ParameterList>
```

changes the smoother for level 2 to be a polynomial smoother.

### 5.2 Parameter validation

By default, MUELU validates the input parameter list. A parameter that is misspelled, is unknown, or has an incorrect value type will cause an exception to be thrown and execution to halt.

☛ Spaces are important within a parameter's name. Please separate words by just one space, and make sure there are no leading or trailing spaces.

The option `print initial parameters` prints the initial list given to the interpreter. The option `print unused parameters` prints the list of unused parameters.

### 5.3 General options

Verbosity level	Description
none	No output
low	Errors, important warnings, and some statistics
medium	Same as low, but with more statistics
high	Errors, all warnings, and all statistics
extreme	Same as high, but also includes output from other packages ( <i>i.e.</i> , ZOLTAN)

**Table 5.1.** Verbosity levels.

Problem type	Multigrid algorithm	Block size	Smoother
unknown	–	–	–
Poisson-2D	Smoothed aggregation	1	Chebyshev
Poisson-3D	Smoothed aggregation	1	Chebyshev
Elasticity-2D	Smoothed aggregation	2	Chebyshev
Elasticity-3D	Smoothed aggregation	3	Chebyshev
ConvectionDiffusion	Petrov-Galerkin AMG	1	Gauss-Seidel
MHD	Unsmoothed aggregation	–	Additive Schwarz method with one level of overlap and ILU(0) as a subdomain solver

**Table 5.2.** Supported problem types (“–” means not set).

<code>problem: type</code>	[string] Type of problem to be solved. Possible values: see Table 5.2. <b>Default:</b> “unknown”.
<code>verbosity</code>	[string] Control of the amount of printed information. Possible values: see Table 5.1. <b>Default:</b> “high”.

number of equations	[int] Number of PDE equations at each grid node. Only constant block size is considered. <b>Default:</b> 1.
max levels	[int] Maximum number of levels in a hierarchy. <b>Default:</b> 10.
cycle type	[string] Multigrid cycle type. Possible values: "V", "W". <b>Default:</b> "V".
problem: symmetric	[bool] Symmetry of a problem. This setting affects the construction of a restrictor. If set to true, the restrictor is set to be the transpose of a prolongator. If set to false, underlying multigrid algorithm makes the decision. <b>Default:</b> true.
xml parameter file	[string] An XML file from which to read additional parameters. In case of a conflict, parameters manually set on the list will override parameters in the file. If the string is empty a file will not be read. <b>Default:</b> "".

## 5.4 Smoothing and coarse solver options

MUELU relies on other TRILINOS packages to provide level smoothers and coarse solvers. IFPACK and IFPACK2 provide standard smoothers (see Table 5.3), and AMESOS and AMESOS2 provide direct solvers (see Table 5.4). Note that it is completely possible to use any level smoother as a direct solver.

MUELU checks parameters `smoother: *` `type` and `coarse: type` to determine:

- what package to use (i.e., is it a smoother or a direct solver);
- (possibly) transform a smoother type
  - IFPACK and IFPACK2 use different smoother type names, e.g., “point relaxation stand-alone” vs “RELAXATION”. MUELU tries to follow IFPACK2 notation for smoother types. Please consult IFPACK2 documentation [11] for more information.

The parameter lists `smoother: *` `params` and `coarse: params` are passed directly to the corresponding package without any examination of their content. Please consult the documentation of the corresponding packages for a list of possible values.

By default, MUELU uses one sweep of symmetric Gauss-Seidel for both pre- and post-smoothing, and SuperLU for coarse system solver.

smoother: type	
RELAXATION	Point relaxation smoothers, including Jacobi, Gauss-Seidel, symmetric Gauss-Seidel, etc. The exact smoother is chosen by specifying relaxation: type parameter in the smoother: params sublist.
Chebyshev	Chebyshev polynomial smoother.
ILUT, RILUK	Local (processor-based) incomplete factorization methods.

**Table 5.3.** Commonly used smoothers provided by IFPACK/IFPACK2. Note that these smoothers can also be used as coarse grid solvers.

coarse: type	AMESOS	AMESOS2	
KLU	x		Default AMESOS solver [7].
KLU2		x	Default AMESOS2 solver [4].
SuperLU	x	x	Third-party serial sparse direct solver [8].
SuperLU_dist	x	x	Third-party parallel sparse direct solver [8].
Umfpack	x		Third-party solver [6].
Mumps	x		Third-party solver [2].

**Table 5.4.** Commonly used direct solvers provided by AMESOS/AMESOS2

smoother: pre or post	[string] Pre- and post-smoother combination. Possible values: "pre" (only pre-smoother), "post" (only post-smoother), "both" (both pre-and post-smoothers), "none" (no smoothing). <b>Default:</b> "both".
smoother: type	[string] Smoother type. Possible values: see Table 5.3. <b>Default:</b> "RELAXATION".
smoother: pre type	[string] Pre-smoother type. Possible values: see Table 5.3. <b>Default:</b> "RELAXATION".

smoother: post type	[string] Post-smoother type. Possible values: see Table 5.3. <b>Default:</b> "RELAXATION".
smoother: params	[ParameterList] Smoother parameters. For standard smoothers, MUELU passes them directly to the appropriate package library.
smoother: pre params	[ParameterList] Pre-smoother parameters. For standard smoothers, MUELU passes them directly to the appropriate package library.
smoother: post params	[ParameterList] Post-smoother parameters. For standard smoothers, MUELU passes them directly to the appropriate package library.
smoother: overlap	[int] Smoother subdomain overlap. <b>Default:</b> 0.
smoother: pre overlap	[int] Pre-smoother subdomain overlap. <b>Default:</b> 0.
smoother: post overlap	[int] Post-smoother subdomain overlap. <b>Default:</b> 0.
coarse: max size	[int] Maximum dimension of a coarse grid. MUELU will stop coarsening once it is achieved. <b>Default:</b> 2000.
coarse: type	[string] Coarse solver. Possible values: see Table ??. <b>Default:</b> "SuperLU".
coarse: params	[ParameterList] Coarse solver parameters. MUELU passes them directly to the appropriate package library.
coarse: overlap	[int] Coarse solver subdomain overlap. <b>Default:</b> 0.

## 5.5 Aggregation options

uncoupled	Attempts to construct aggregates of optimal size ( $3^d$ nodes in $d$ dimensions). Each process works independently, and aggregates cannot span multiple processes.
coupled	Attempts to construct aggregates of optimal size ( $3^d$ nodes in $d$ dimensions). Aggregates are allowed to cross processor boundaries. <b>Use carefully.</b> If unsure, use uncoupled instead.
brick	Attempts to construct rectangular aggregates

**Table 5.5.** Available coarsening schemes.

aggregation: type	[string] Aggregation scheme. Possible values: see Table 5.5. <b>Default:</b> "uncoupled".
aggregation: ordering	[string] Node ordering strategy. Possible values: "natural" (local index order), "graph" (filtered graph breadth-first order), "random" (random local index order). <b>Default:</b> "natural".
aggregation: drop scheme	[string] Connectivity dropping scheme for a graph used in aggregation. Possible values: "classical", "distance laplacian". <b>Default:</b> "classical".
aggregation: drop tol	[double] Connectivity dropping threshold for a graph used in aggregation. <b>Default:</b> 0.0.
aggregation: min agg size	[int] Minimum size of an aggregate. <b>Default:</b> 2.
aggregation: max agg size	[int] Maximum size of an aggregate (-1 means unlimited). <b>Default:</b> -1.
aggregation: brick x size	[int] Number of points for x axis in "brick" aggregation (limited to 3). <b>Default:</b> 2.



aggregation: brick y size	[int] Number of points for y axis in "brick" aggregation (limited to 3). <b>Default:</b> 2.
aggregation: brick z size	[int] Number of points for z axis in "brick" aggregation (limited to 3). <b>Default:</b> 2.
aggregation: Dirichlet threshold	[double] Threshold for determining whether entries are zero during Dirichlet row detection. <b>Default:</b> 0.0.
aggregation: export visualization data	[bool] Export data for visualization post-processing. <b>Default:</b> false.
aggregation: output filename	[string] Filename to write VTK visualization data to. <b>Default:</b> "".
aggregation: output file: time step	[int] Time step ID for non-linear problems. <b>Default:</b> 0.
aggregation: output file: iter	[int] Iteration for non-linear problems. <b>Default:</b> 0.
aggregation: output file: agg style	[string] Style of aggregate visualization. <b>Default:</b> Point Cloud.
aggregation: output file: fine graph edges	[bool] Whether to draw all fine node connections along with the aggregates. <b>Default:</b> false.
aggregation: output file: coarse graph edges	[bool] Whether to draw all coarse node connections along with the aggregates. <b>Default:</b> false.
aggregation: output file: build colormap	[bool] Whether to output a random colormap in a separate XML file. <b>Default:</b> false.

## 5.6 Rebalancing options

repartition: enable	[bool] Rebalancing on/off switch. <b>Default:</b> false.
---------------------	--

repartition: partitioner	[string] Partitioning package to use. Possible values: "zoltan" (ZOLTAN library), "zoltan2" (ZOLTAN2 library). <b>Default:</b> "zoltan2".
repartition: params	[ParameterList] Partitioner parameters. MUELU passes them directly to the appropriate package library.
repartition: start level	[int] Minimum level to run partitioner. MUELU does not rebalance levels finer than this one. <b>Default:</b> 2.
repartition: min rows per proc	[int] Minimum number of rows per processor. If the actual number is smaller, then rebalancing occurs. <b>Default:</b> 800.
repartition: max imbalance	[double] Maximum nonzero imbalance ratio. If the actual number is larger, the rebalancing occurs. <b>Default:</b> 1.2.
repartition: remap parts	[bool] Postprocessing for partitioning to reduce data migration. <b>Default:</b> true.
repartition: rebalance P and R	[bool] Explicit rebalancing of R and P during the setup. This speeds up the solve, but slows down the setup phases. <b>Default:</b> false.

## 5.7 Multigrid algorithm options

multigrid algorithm	[string] Multigrid method. Possible values: see Table 5.6. <b>Default:</b> "sa".
sa: damping factor	[double] Damping factor for smoothed aggregation. <b>Default:</b> 1.33.

sa	Classic smoothed aggregation [15]
unsmoothed	Aggregation-based, same as sa but without damped Jacobi prolongator improvement step
pg	Prolongator smoothing using $A$ , restriction smoothing using $A^T$ , local damping factors [12]
emin	Constrained minimization of energy in basis functions of grid transfer operator [18, 10]

**Table 5.6.** Available multigrid algorithms for generating grid transfer matrices.

sa: use filtered matrix	[bool] Matrix to use for smoothing the tentative prolongator. The two options are: to use the original matrix, and to use the filtered matrix with filtering based on filtered graph used for aggregation. <b>Default:</b> true.
filtered matrix: use lumping	[bool] Lump (add to diagonal) dropped entries during the construction of a filtered matrix. This allows user to preserve constant nullspace. <b>Default:</b> true.
filtered matrix: reuse eigenvalue	[bool] Skip eigenvalue calculation during the construction of a filtered matrix by reusing eigenvalue estimate from the original matrix. This allows us to skip heavy computation, but may lead to poorer convergence. <b>Default:</b> true.
emin: iterative method	[string] Iterative method to use for energy minimization of initial prolongator in energy-minimization. Possible values: "cg" (conjugate gradient), "gmres" (generalized minimum residual), "sd" (steepest descent). <b>Default:</b> "cg".
emin: num iterations	[int] Number of iterations to minimize initial prolongator energy in energy-minimization. <b>Default:</b> 2.
emin: num reuse iterations	[int] Number of iterations to minimize the reused prolongator energy in energy-minimization. <b>Default:</b> 1.

emin: pattern	[string] Sparsity pattern to use for energy minimization. Possible values: "AkPtent". <b>Default:</b> "AkPtent".
emin: pattern order	[int] Matrix order for the "AkPtent" pattern. <b>Default:</b> 1.

## 5.8 Reuse options

Reuse options are a way for a user to speed up the setup stage of multigrid. The main requirement to use reuse is that the matrix' graph structure does not change. Only matrix values are allowed to change.

The reuse options control the degree to which multigrid hierarchy is preserved for a subsequent setup call.

In addition, please note that not all combinations of multigrid algorithms and reuse options are valid, or even make sense. For instance, the "emin" reuse option should only be use with "emin" multigrid algorithm.

Table 5.7 contains the information about different reuse options. The options are ordered in increasing number of reuse components, from the no reuse to the full reuse ("full").

none	No reuse
S	Reuse only the symbolic information of the level smoothers.
tP	Reuse tentative prolongator. The graphs of smoothed prolongator and matrices in Galerkin product are reused only if filtering is not being used ( <i>i.e.</i> , either <code>sa: use filtered matrix</code> or <code>aggregation: drop tol</code> is false)
emin	Reuse old prolongator as an initial guess to energy minimization, and reuse the prolongator pattern
RP	Reuse smoothed prolongator and restrictor. Smoothers are recomputed. ☛ RP should reuse matrix graphs for matrix-matrix product, but currently that is disabled as only EPETRA supports it.
RAP	Recompute only the finest level smoothers, reuse all other operators
full	Reuse everything

**Table 5.7.** Available coarsening schemes.

<code>reuse: type</code>	[string] Reuse options for consecutive hierarchy construction. This speeds up the setup phase, but may lead to poorer convergence. Possible values: see Table 5.7. <b>Default:</b> "none".
--------------------------	--

## 5.9 Miscellaneous options

<code>export data</code>	[ParameterList] Exporting a subset of the hierarchy data in a file. Currently, the list can contain any of three parameter names ("A", "P", "R") of type string and value "{levels separated by commas}". A matrix with a name "X" is saved in three MatrixMarket files: a) data is saved in <i>X_level.mm</i> ; b) its row map is saved in <i>rowmap_X_level.mm</i> ; c) its column map is saved in <i>colmap_X_level.mm</i> .
<code>print initial parameters</code>	[bool] Print parameters provided for a hierarchy construction. <b>Default:</b> true.
<code>print unused parameters</code>	[bool] Print parameters unused during a hierarchy construction. <b>Default:</b> true.
<code>transpose: use implicit</code>	[bool] Use implicit transpose for the restriction operator. <b>Default:</b> false.



# Chapter 6

## MUEMEX: The MATLAB Interface for MUEL

MUEMEX is MUEL's interface to the MATLAB environment. It allows access to a limited set of routines either MUEL as a preconditioner, Belos as a solver and Epetra or Tpetra for data structures. It is designed to provide access to MUEL's aggregation and solver routines from MATLAB and does little else. MUEMEX allows users to setup and solve arbitrarily many problems, so long as memory suffices. More than one problem can be set up simultaneously.

### 6.1 Cmake Configure and Make

To use MUEMEX, Trilinos must be configured with (at least) the following options:

```
export TRILINOS_HOME=/path/to/your/Trilinos/source/directory
cmake \
  -D Trilinos_ENABLE_EXPLICIT_INSTANTIATION:BOOL=ON \
  -D Trilinos_ENABLE_Amesos:BOOL=ON \
  -D Trilinos_ENABLE_Amesos2:BOOL=ON \
  -D Amesos2_ENABLE_KLU2:BOOL=ON \
  -D Trilinos_ENABLE_Aztec00:BOOL=ON \
  -D Trilinos_ENABLE_Epetra:BOOL=ON \
  -D Trilinos_ENABLE_EpetraExt:BOOL=ON \
  -D Trilinos_ENABLE_Fortran:BOOL=OFF \
  -D Trilinos_ENABLE_Ipack:BOOL=ON \
  -D Trilinos_ENABLE_Ipack2:BOOL=ON \
  -D Trilinos_ENABLE_MueLu:BOOL=ON \
  -D Trilinos_ENABLE_Teuchos:BOOL=ON \
  -D Trilinos_ENABLE_Tpetra:BOOL=ON \
  -D TPL_ENABLE_MPI:BOOL=OFF \
  -D TPL_ENABLE_MATLAB:BOOL=ON \
  -D MATLAB_ROOT:STRING="<>my matlab root"> \
  -D MATLAB_ARCH:STRING="<>my matlab os string"> \
  -D Trilinos_EXTRA_LINK_FLAGS="-lrt -lm -lgfortran" \
  ${TRILINOS_HOME}
```

Since MUEMEX supports both the Epetra and Tpetra linear algebra libraries, you have to have both enabled in order to build MUEMEX.

☛ If you turn off either Epetra or Tpetra then you will run into an error message: *MueMex requires Epetra, Tpetra and MATLAB*.

Most additional options can be specified as well. It is important to note that MUEMEX does not work properly with MPI, hence MPI must be disabled in order to compile MUEMEX. The MATLAB\_ARCH option is new to the cmake build system, and involves the MATLAB-specific architecture code for your system. There is currently no automatic way to extract this, so it must be user-specified. As of MATLAB 7.9 (R2009b), common arch codes are:

Code	OS
glnx86	32-bit Linux (intel/amd)
glnxa64	64-bit Linux (intel/amd)
maci64	64-bit MacOS
maci	32-bit MacOS

On 64-bit Intel/AMD architectures, Trilinos and all relevant TPLs (note: this includes BLAS and LAPACK) must be compiled with the `-fPIC` option. This necessitates adding:

```
-D CMAKE_CXX_FLAGS:STRING="-fPIC" \  
-D CMAKE_C_FLAGS:STRING="-fPIC" \  
-D CMAKE_Fortran_FLAGS:STRING="-fPIC" \
```

to the cmake configure line.

The additional linker flags specified in `Trilinos_EXTRA_LINK_FLAGS` may slightly vary depending on the system and the exact configuration. But the given parameters may work for most Linux based systems. If you encounter an error message like *Target "mudemex.mexa64" links to item "-Wl,-rpath-link,/opt/matlab/bin/glnxa64" which has leading or trailing whitespace.* you have to add some options to the `Trilinos_EXTRA_LINK_FLAGS` variable. At least adding `-lm` should be safe and fix the error message.

### 6.1.1 BLAS & LAPACK Option #1: Static Builds

Trilinos does not play nicely with MATLAB's default LAPACK and BLAS on 64-bit machines. If MUEMEX randomly crashes when you run with any Krylov method that has orthogonalization, chances are MUEMEX is finding the wrong BLAS/LAPACK libraries. This leaves you with one of two options. The first is to build them both *statically* and then specify them as follows:

```
-D LAPACK_LIBRARY_DIRS:STRING="<path to my lapack.a>" \  
-D BLAS_LIBRARY_DIRS:STRING="<path to my blas.a>" \
```



Using static linking for LAPACK and BLAS prevents MATLAB's default libraries to take precedence.

### 6.1.2 BLAS & LAPACK Option #2: LD\_PRELOAD

The second option is to use LD\_PRELOAD to tell MATLAB exactly which libraries to use. For this option, you can use the dynamic libraries installed on your system. Before starting MATLAB, set LD\_PRELOAD to the paths of libstdc++.so corresponding to the version of GCC used to build Trilinos, and the paths of libblas.so and liblapack.so on your local system.

For example, if you use bash, you'd do something like this

```
export LD_PRELOAD=<path>/libstdc++.so:<path>/libblas.so:<path>/liblapack.so
```

For csh / tcsh, do this

```
setenv LD_PRELOAD <path>/libstdc++.so:<path>/libblas.so:<path>/liblapack.so
```

### 6.1.3 Running MATLAB

Before you run MATLAB you have to make sure that MATLAB is using the same libraries that have been used for compiling MUEMEX. This includes the libstdc++.so and depending whether you turned on/off fortran also libgfortran.so. Please make sure that the correct libraries and paths are declared in the LD\_PRELOAD variable. You can refer to section 6.1.2 to see how the LD\_PRELOAD variable is set.

For a 64 bit Linux system using the bash the command should look like

```
export LD_PRELOAD=/usr/lib64/libstdc++.so.6:/usr/lib64/libgfortran.so.3:  
$LD_PRELOAD
```

to add the libstdc++.so and libgfortran.so to the existing LD\_PRELOAD variable. Then run the MATLAB executable in the same shell window.

☛ Note, that this step is necessary even if you statically linked BLAS and LAPACK.

If you are unsure which libraries have to be set in the LD\_PRELOAD variable you will find out latest if you start MATLAB and try to run MUEMEX. It will throw some error messages with the missing library names. For a 64 bit Linux system the standard libraries usually can be found in /usr/lib64 or /usr/lib (for a 32 bit system).

## 6.2 Using MUEMEX

MUEMEX is designed to be interfaced with via the MATLAB script `muelu.m`. There are five modes in which MUEMEX can be run:

1. **Setup Mode** — Performs the problem setup for MUELU. Depending on whether or not the Linear Algebra option is used, MUEMEX creates either an unpreconditioned Epetra problem, an Epetra problem with MUELU, or a Tpetra problem with MUELU. The default is `tpetra`. The `epetra` mode only supports real-valued matrices, while `tpetra` supports both real and complex and will infer the scalar type from the matrix passed during setup. This call returns a problem handle used to reference the problem in the future, and (optionally) the operator complexity, if a preconditioner is being used.
2. **Solve Mode** — Given a problem handle and a right-hand side, MUEMEX solves the problem specified. Setup mode must be called before solve mode.
3. **Cleanup Mode** — Frees the memory allocated to internal MUELU, Epetra and Tpetra objects. This can be called with a particular problem handle, in which case it frees that problem, or without one, in which case all MUEMEX memory is freed.
4. **Status Mode** — Prints out status information on problems which have been set up. Like cleanup, it can be called with or without a particular problem handle.
5. **Get Mode** — Get information from a MueLu hierarchy that has been generated. Given the problem handle, a level number and the name of the field, returns the appropriate array or scalar as a MATLAB object.

All of these modes, with the exception of status and cleanup take option lists which will be directly converted into `Teuchos::ParameterList` objects by MUEMEX, as key-value pairs. Options passed during setup will apply to the MUELU preconditioner, and options passed during a solve will apply to Belos.

### 6.2.1 Setup Mode

Setup mode is called as follows:

```
>> [h, oc] = muelu('setup', A[, 'parameter', value,...])
```

The parameter `A` represents the sparse matrix to perform aggregation on and the parameter/value pairs represent standard MUELU options.

The routine returns a problem handle, `h`, and the operator complexity `oc` for the operator. In addition to the standard options, setup mode has one unique option of its own:

### 6.2.2 Solve Mode

Solve mode is called as follows:

```
>> [x, its] = muelu(h[, A], b[, 'parameter', value,...])
```

The parameter `h` is a problem handle returned by the setup mode call, `A` is the sparse matrix with which to solve and `b` is the right-hand side. Parameter/value pairs to configure the Belos solver are listed as above. If `A` is not supplied, the matrix provided when setting up the problem will be used. `x` is the solution multivector with the same dimensions as `b`, and `its` is the number of iterations Belos needed to solve the problem.

All of these options are taken directly from Belos, so consult its manual for more information. Belos output style and verbosity settings are implemented as enums, but can be set as strings in MUEMEX. For example:

```
>> x = muelu(0, b, 'Verbosity', 'Warnings + IterationDetails', ...
               'Output Style', 'Brief');
```

Verbosity settings can be separated by spaces, '+' or ','. Belos::Brief is the default output style.

### 6.2.3 Cleanup Mode

Cleanup mode is called as follows:

```
>> muelu('cleanup', [, h])
```

The parameter `h` is a problem handle returned by the setup mode call and is optional. If `h` is provided, that problem is cleaned up. If the option is not provided all currently set up problems are cleaned up.

### 6.2.4 Status Mode

Status mode is called as follows:

```
>> muelu('status', [, h])
```

The parameter `h` is a problem handle returned by the setup mode call and is optional. If `h` is provided, status information for that problem is printed. If the option is not provided all currently set up problems have status information printed.

### 6.2.5 Get Mode

Get mode is called as follows:

```
>> muelu('get', h, level, fieldName[, typeHint])
```

The parameter `h` is the problem handle, and `level` is an integer that identifies the level within the hierarchy containing the desired data. `fieldName` is a string that identifies the field within the level, e.g. 'Nullspace'. `typeHint` is an optional parameter that tells MueMex what data type to expect from the level. This is a string, with possible values 'matrix', 'multivector', 'lovector' (ordinal vector), or 'scalar'. MueMex will attempt to guess the type from `fieldName` but `typeHint` may be required.

### 6.2.6 Tips and Tricks

Internally, MATLAB represents all data as doubles unless you go through efforts to do otherwise. MUEMEX detects integer parameters by a relative error test, seeing if the relative difference between the value from MATLAB and the value of the `int-typecast` value are less than  $1e-15$ . Unfortunately, this means that MUEMEX will choose the incorrect type for parameters which are doubles that happen to have an integer value (a good example of where this might happen would be the parameter 'smoother Chebyshev: alpha', which defaults to 30.0). Since MUEMEX does no internal typechecking of parameters (it uses MUELU's internal checks), it has no way of detecting this conflict. From the user's perspective, avoiding this is as simple as adding a small perturbation (greater than a relative  $1e-15$ ) to the parameter that makes it non-integer valued.

# References

- [1] Trilinos CMake Quickstart. [http://trilinos.org/build\\_instructions.html](http://trilinos.org/build_instructions.html), 2014.
- [2] Patrick R. Amestoy, Iain S. Duff, Jean-Yves L’Excellent, and Jacko Koster. MUMPS: a general purpose distributed memory sparse solver. In *Applied Parallel Computing. New Paradigms for HPC in Industry and Academia*, pages 121–130. Springer, 2001.
- [3] Roscoe A. Bartlett. Teuchos::RCP beginners guide. Technical Report SAND2004-3268, Sandia National Labs, 2010.
- [4] Eric Bavier, Mark Hoemmen, Sivasankaran Rajamanickam, and Heidi Thornquist. Amesos2 and Belos: Direct and iterative solvers for large sparse linear systems. *Scientific Programming*, 20(3):241–255, 2012.
- [5] William L. Briggs, Van Emden Henson, and Steve F. McCormick. *A multigrid tutorial*. SIAM, 2nd edition, 2000.
- [6] Timothy A. Davis. Algorithm 832: UMFPACK v4.3 — an unsymmetric-pattern multifrontal method. *ACM Transactions on Mathematical Software (TOMS)*, 30(2):196–199, 2004.
- [7] Timothy A. Davis and Ekanathan Palamadai Natarajan. Algorithm 907: KLU, a direct sparse solver for circuit simulation problems. *ACM Transactions on Mathematical Software (TOMS)*, 37(3):36, 2010.
- [8] Xiaoye S. Li, James W. Demmel, John R. Gilbert, Laura Grigori, Meiyue Shao, and Ichitaro Yamazaki. SuperLU Users’ Guide. 2011.
- [9] Paul Lin, Matthew Bettencourt, Stefan Domino, Travis Fisher, Mark Hoemmen, Jonathan J. Hu, Eric Phipps, Andrey Prokopenko, Sivasankaran Rajamanickam, Christopher Siefert, and Stephen Kennon. Towards extreme-scale simulations for low Mach fluids with second-generation Trilinos. *Parallel Processing Letters*, 24(04):1442005, 2014.
- [10] L. Olson, J. Schroder, and R. Tuminaro. A general interpolation strategy for algebraic multigrid using energy minimization. *SIAM Journal on Scientific Computing*, 33(2):966–991, 2011.
- [11] Andrey Prokopenko, Christopher M. Siefert, Jonathan J. Hu, Mark Hoemmen, and Alicia Klinvex. Ifpack2 Users Guide 1.0. Technical Report SAND2016-5338, Sandia National Labs, 2016.
- [12] Marzio Sala and Raymond S. Tuminaro. A new Petrov-Galerkin smoothed aggregation preconditioner for nonsymmetric linear systems. *SIAM Journal on Scientific Computing*, 31(1):143–166, 2008.

- [13] Heidi Thornquist, Roscoe A. Bartlett, Mark Hoemmen, Christopher Baker, and Michael Heroux. Teuchos: Trilinos tools library. <http://trilinos.org/packages/teuchos>, 2014.
- [14] Ulrich Trottenberg, Cornelis Oosterlee, and Anton Schüller. *Multigrid*. Elsevier Academic Press, 2001.
- [15] P. Vaněk, J. Mandel, and M. Brezina. Algebraic multigrid based on smoothed aggregation for second and fourth order problems. *Computing*, 56:179–196, 1996.
- [16] Tobias A. Wiesner. *Flexible Aggregation-based Algebraic Multigrid Methods for Contact and Flow Problems*. PhD thesis, 2014.
- [17] Tobias A. Wiesner, Michael W. Gee, Andrey Prokopenko, and Jonathan J. Hu. The MueLu tutorial. <http://trilinos.org/packages/muelu/muelu-tutorial>, 2014. SAND2014-18624R.
- [18] Tobias A. Wiesner, Raymond S. Tuminaro, Wolfgang A. Wall, and Michael W. Gee. Multigrid transfers for nonsymmetric systems based on Schur complements and Galerkin projections. *Numerical Linear Algebra with Applications*, 21(3):415–438, 2014.

# Appendix A

## Copyright and License

MueLu: A package for multigrid based preconditioning

Copyright 2012 Sandia Corporation

Under the terms of Contract DE-AC04-94AL85000 with Sandia Corporation, the U.S. Government retains certain rights in this software.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

1. Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.
3. Neither the name of the Corporation nor the names of the contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY SANDIA CORPORATION "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL SANDIA CORPORATION OR THE CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.





# Appendix B

## ML compatibility

MUELU provides a basic compatibility layer for ML parameter lists. This allows ML users to quickly perform some experiments with MUELU.

**First and most important:** Long term, we would like to have users use the new MUELU interface, as that is where most of new features will be made accessible. One should make note of the fact that it may not be possible to make ML deck do exactly same things in ML and MUELU, as internally ML implicitly makes some decision that we have no control over and which could be different from MUELU.

There are basically two distinct ways to use ML input parameters with MUELU:

**MLParameterListInterpreter:** This class is the pendant of the `ParameterListInterpreter` class for the MUELU parameters. It accepts parameter lists or XML files with ML parameters and generates a MUELU multigrid hierarchy. It supports only a well-defined subset of ML parameters which have a equivalent parameter in MUELU.

**ML2MueLuParameterTranslator:** This class is a simple wrapper class which translates ML parameters to the corresponding MUELU parameters. It has to be used in combination with the MUELU `ParameterListInterpreter` class to generate a MUELU multigrid hierarchy. It is also meant to be used in combination with the `CreateEpetraPreconditioner` and `CreateTpetraPreconditioner` routines (see §3.3). It supports only a small subset of the ML parameters.

## B.1 Usage of ML parameter lists with MUELU

### B.1.1 MLParameterListInterpreter

The `MLParameterListInterpreter` directly accepts a `ParameterList` containing ML parameters. It also interprets the null space: vectors and the null space: dimension ML parameters. However, it is recommended to provide the near null space vectors directly in the MUELU way as shown in the following code snippet.

```

Teuchos::RCP<Tpetra::CrsMatrix<> > A;
// create A here ...

// XML file containing ML parameters
std::string xmlFile = "mlParameters.xml"
Teuchos::ParameterList paramList;
Teuchos::updateParametersFromXmlFileAndBroadcast(xmlFile, Teuchos::Ptr<Teuchos
::ParameterList>(&paramList), *comm);

// use ParameterListInterpreter with MueLu parameters as input
Teuchos::RCP<HierarchyManager> mueluFactory = Teuchos::rcp(new
    MLParameterListInterpreter(*paramList));

RCP<Hierarchy> H = mueluFactory->CreateHierarchy();
H->GetLevel(0)->Set<RCP<Matrix> >("A", A);
H->GetLevel(0)->Set("Nullspace", nullspace);
H->GetLevel(0)->Set("Coordinates", coordinates);
mueluFactory->SetupHierarchy(*H);

```

Note that the `MLParameterListInterpreter` only supports a basic set of ML parameters allowing to build smoothed aggregation transfer operators (see §B.2 for a list of compatible ML parameters).

### B.1.2 ML2MueLuParameterTranslator

The `ML2MueLuParameterTranslator` class is a simple wrapper translating ML parameters to the corresponding MUELU parameters. This allows the usage of the simple `CreateEpetraPreconditioner` and `CreateTpetraPreconditioner` interface with ML parameters:

```

Teuchos::RCP<Tpetra::CrsMatrix<> > A;
// create A here ...

// XML file containing ML parameters
std::string xmlFile = "mlParameters.xml"
Teuchos::ParameterList paramList;
Teuchos::updateParametersFromXmlFileAndBroadcast(xmlFile, Teuchos::Ptr<Teuchos
::ParameterList>(&paramList), *comm);

// translate ML parameters to MueLu parameters
RCP<ParameterList> mueluParamList = Teuchos::getParametersFromXmlString(MueLu
::ML2MueLuParameterTranslator::translate(paramList, "SA"));

Teuchos::RCP<MueLu::TpetraOperator> mueluPreconditioner =
    MueLu::CreateTpetraPreconditioner(A, mueluParamList);

```

In a similar way, ML input parameters can be used with the standard MUELU parameter list interpreter class. Note that the near null space vectors have to be provided in the MUELU way.

```
Teuchos::RCP<Tpetra::CrsMatrix<> > A;
// create A here ...

// XML file containing ML parameters
std::string xmlFile = "mlParameters.xml"
Teuchos::ParameterList paramList;
Teuchos::updateParametersFromXmlFileAndBroadcast(xmlFile, Teuchos::Ptr<Teuchos
::ParameterList>(&paramList), *comm);

// translate ML parameters to MueLu parameters
RCP<ParameterList> mueluParamList = Teuchos::getParametersFromXmlString(MueLu
::ML2MueLuParameterTranslator::translate(paramList, "SA"));

// use ParameterListInterpreter with MueLu parameters as input
Teuchos::RCP<HierarchyManager> mueluFactory = Teuchos::rcp(new
ParameterListInterpreter(*mueluParamList));

RCP<Hierarchy> H = mueluFactory->CreateHierarchy();
H->GetLevel(0)->Set<RCP<Matrix> >("A", A);
H->GetLevel(0)->Set("Nullspace", nullspace);
H->GetLevel(0)->Set("Coordinates", coordinates);
mueluFactory->SetupHierarchy(*H);
```

Note that the set of supported ML parameters is very limited. Please refer to §B.2 for a list of all compatible ML parameters.

## B.2 Compatible ML parameters

### B.2.1 General ML options

ML output	[int] Control of the amount of printed information. Possible values: 0-10 with 0=no output and 10=maximum verbosity. <b>Default:</b> 0 <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.
PDE equations	[int] Number of PDE equations at each grid node. Only constant block size is considered. <b>Default:</b> 1 <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.

max levels	[int] Maximum number of levels in a hierarchy. <b>Default:</b> 10 <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.
prec type	[string] Multigrid cycle type. Possible values: "MGV", "MGW". Other values are NOT supported by MueLu. <b>Default:</b> "MGV" <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.

## B.2.2 Smoothing and coarse solver options

smoother: type	[string] Smoother type for fine- and intermedium multigrid levels. Possible values: "Jacobi", "Gauss-Seidel", "symmetric Gauss-Seidel", "Chebyshev", "ILU". <b>Default:</b> "symmetric Gauss-Seidel" <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.
smoother: sweeps	[int] Number of smoother sweeps for relaxation based level smoothers. In case of Chebyshev smoother it denotes the polynomial degree. <b>Default:</b> 2 <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.
smoother: damping factor	[double] Damping factor for relaxation based level smoothers. <b>Default:</b> 1.0 <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.
smoother: Chebyshev alpha	[double] Eigenvalue ratio for Chebyshev level smoother. <b>Default:</b> 20 <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.
smoother: pre or post	[string] Pre- and post-smoother combination. Possible values: "pre" (only pre-smoother), "post" (only post-smoother), "both" (both pre-and post-smoothers). <b>Default:</b> "both" <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.

max size	[int] Maximum dimension of a coarse grid. ML will stop coarsening once it is achieved. <b>Default:</b> 128 <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.
coarse: type	[string] Solver for coarsest level. Possible values: "Amesos-KLU", "Amesos-Superlu" (depending on MUELU installation). <b>Default:</b> "Amesos-KLU" <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.

### B.2.3 Transfer operator options

energy minimization: enable	[int] Enable energy minimization transfer operators (using Petrov-Galerkin approach). <b>Default:</b> 0 <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.
aggregation: damping factor	[double] Damping factor for smoothed aggregation. <b>Default:</b> 1.33 <b>Compatibility:</b> MLParameterListInterpreter, ML2MueLuParameterTranslator.

### B.2.4 Rebalancing options

repartition: enable	[int] Rebalancing on/off switch. Only limited support for repartitioning. Does not use provided node coordinates. <b>Default:</b> 0 <b>Compatibility:</b> MLParameterListInterpreter.
repartition: start level	[int] Minimum level to run partitioner. MUELU does not rebalance levels finer than this one. <b>Default:</b> 1 <b>Compatibility:</b> MLParameterListInterpreter.
repartition: min per proc	[int] Minimum number of rows per processor. If the actual number is smaller, then rebalancing occurs. <b>Default:</b> 512 <b>Compatibility:</b> MLParameterListInterpreter.

repartition: max min  
ratio

[double] Maximum nonzero imbalance ratio. If the actual number is larger, the rebalancing occurs. **Default:** 1.3 **Compatibility:** MLParameterListInterpreter.

## DISTRIBUTION:

- 1 Tobias Wiesner  
Institute for Computational Mechanics  
Technische Universität München  
Boltzmanstraße 15  
85747 Garching, Germany
- 1 MS 1320 Michael Heroux, 1446
- 1 MS 1318 Robert Hoekstra, 1446
- 1 MS 1320 Mark Hoemmen, 1446
- 1 MS 1320 Paul Lin, 1446
- 1 MS 1318 Andrey Prokopenko, 1426
- 1 MS 1322 Christopher Siefert, 1443
- 1 MS 0899 Technical Library, 9536 (electronic copy)

## DISTRIBUTION:

1	MS 9159	Jonathan Hu, 1426
1	MS 9159	Paul Tsuji, 1442
1	MS 9159	Raymond Tuminaro, 1442
1	MS 0899	Technical Library, 8944 (electronic copy)





